

Secure Communication Protocols

Abstract

In this assignment, we implement a secure communication protocol consisting of three consecutive steps. First, we use a Diffie-Hellman^{1,2} key exchange implementation to craft a shared private key between two actors using a 2048-bit cyclic group defined by the IETF. We then use an implementation of the Blum Blum Shub³ algorithm to generate a cryptographically strong pseudo-random shared private key, using the previously exchanged key as seed. Finally, we use the generated random key to encrypt and decrypt messages between the two actors using AES⁴ ciphers in ECB mode. For each step, we assess the vulnerabilities of our protocol and propose more secure alternatives.

1 Introduction

Our main goal is to be able to communicate confidential data over an insecure communication channel. To do so, we have to encrypt our data before sending it, and decrypt it after receiving it, so that it only appears encrypted on the communication channel. To perform these encryptions, we want to use a symmetric cipher allowing us to efficiently compute ciphertext from plaintext and vice versa. Before using symmetric encryption, we have to solve the key exchange problem, because we still need to communicate the private key used by the two actors over an insecure channel. One answer to this first problem is to use public-key cryptography, and more specifically the Diffie-Hellman key exchange protocol.

1.1 Diffie-Hellman

This key exchange scheme allows us to create a shared private key by only communicating public data. It relies on the computational difficulty of computing discrete logarithms. It has been formulated in 1976 by Whitfield Diffie and Martin Hellman^{1,2}. We use it to generate a first shared private key between two actors, Alice and Bob.

1.2 Blum Blum Shub

To increase the strength of the exchanged key, we use it as a seed to a cryptographically strong pseudo-random number generator. The chosen stream generator is Blum Blum Shub³, which outputs the least significant bit of x_{n+1} at each step of the following sequence.

$$x_{n+1} = x_n^2 \bmod M \quad \text{with } x_0 = \text{seed}, M \text{ the product of two large primes } p \text{ and } q, \text{ with } p \equiv 3 \bmod 4 \text{ and } q \equiv \bmod 4$$

1.3 Advanced Encryption Standard

To encrypt their messages, Alice and Bob can now use a symmetric cipher with their randomized shared private key. Here, we use AES in Electronic Code Book mode. However, now that Alice and Bob exchanged their keys, they could use any symmetric cipher and mode of operation.

Using these cryptographic primitives, we implement a complete communication protocol allowing Alice and Bob to communicate confidential data securely.

2 Design and Implementation

2.1 Diffie-Hellman

The python implementation of the Diffie-Hellman key exchange is fairly straightforward. There are two main steps in this scheme : creating a public key from our private key, and combining the other party's public key with our private key to generate the shared private key. These two steps are represented by the two main modes of operation of `keygen.py` : generate and merge, to be passed to the `--mode` argument. When running the script in `--mode generate`, the main function used is the following.

```
def pubkeygen(prime, root, secret):  
    assert(is_prime(prime))  
    assert(is_primitive_root(root, prime))  
  
    return pow(root, secret, prime)
```

And when used in `--mode merge`, the main computation is done in this function.

```
def shared_secret_key(secret, other_public_key, prime):  
    return pow(other_public_key, secret, prime)
```

Note that in both functions, we use the `pow()` method with three arguments rather than exponentiating and then taking the modulo. This is because on large numbers, the `pow()` method is significantly faster at computing the modulo than the raw computation of exponentiation followed by division. This is demonstrated in Figure 2. This program is designed to accept decimal or hexadecimal representation of integers for the `--secret`, `--prime`, `--root`, and `--public` arguments.

The rest of the script is mostly dedicated to argument parsing, parameters loading and tests. The tests are run when using the script in `--mode test`. The test data comes from the RFC 5114⁵ memo which describes standard Diffie-Hellman groups and their associated test data. Results of the tests can be found in Table 1. The default group used by the script if no `--prime` and `--root` are passed is the 2048-bit MODP Group with 256-bit Prime Order Subgroup which can be found in the memo.

Further specification for the other command line arguments can be found in the `README.pdf` file.

2.2 Blum Blum Shub

The main computation of the BBS algorithm is also quite simple, as bits are generated incrementally by taking the least significant bit of each modular exponentiation. The generated bits are then merged together and interpreted as a binary number.

```
def generate_random(seed, size):  
  
    bits = []  
  
    for _ in range(size):  
        seed = pow(seed, 2, M)  
        bits.append(bin(seed)[-1])  
  
    return int(''.join(bits), 2)
```

Although testing for statistical randomness usually involves a series of tests⁶, we can visually check for pseudo-randomness by looking at the bitmap created by the generated bits. If we do not identify any regular pattern, we can assume that the randomness is reasonable for the key to be secure enough. This bitmap test can be seen in Figure 1.

2.3 Advanced Encryption Standard

The implementation of a symmetric cipher is not part of this assignment as it has already been done in Assignment 1. We choose AES for its ingenuity and subsequent popularity, and we use the Electronic Code Book mode of operation so that we do not have to transfer initialization vectors or nonces with the message. In a real-world application, we would prefer to use another mode as ECB is semantically insecure. Given a plaintext, ECB will always produce the same ciphertext each time. For a small amount of short text messages, this mode of operation is secure enough for our protocol.

```
def encrypt(plaintext, key):
    cipher = AES.new(key, AES.MODE_ECB)
    if len(plaintext) % 16 != 0:
        plaintext = pad(plaintext, 16)
    ciphertext = cipher.encrypt(plaintext)
    return ciphertext
```

For encryption, we have to pad the plaintext if it does not fit the block size (16 bytes, 128 bits) as AES is a block cipher. During the decryption process, we unpad the decrypted ciphertext in the similar way.

```
def decrypt(ciphertext, key):
    cipher = AES.new(key, AES.MODE_ECB)
    plaintext = cipher.decrypt(ciphertext)
    if len(plaintext) % 16 != 0:
        plaintext = unpad(plaintext, 16)
    return plaintext
```

Although most of the work here is done by an external library⁷, we can still test the global implementation by using NIST test data⁸. This is done in Table 2.

2.4 Putting it all together

Now that we have seen how the different primitives are implemented, we can combine them to create the secure communication protocol we were seeking. To demonstrate how this can be done, a shell script describing a use case is given. The file **demo.sh** serves as an example of how these primitives can be used in a modular way. Users may chose to run the programs quietly, or in verbose mode to display the parameters used. They may chose to write the generated keys and messages to files. They may chose to change the parameters of the different primitives. This demo shell script is POSIX-compliant and should run on any reasonable shell.

The other shell script provided, **test.sh**, allows to run the implementation tests of the different primitives.

3 Test Results

Table 1 shows the equivalence of test data from RFC 5114⁵ and keys generated by our own implementation. The test data can be found in the **files/2048-bit MODP Group** folder.

Figure 2 shows the speed comparison of modular exponentiation methods. We observe exponential time complexity as the exponent gets larger using naive computation, whereas **pow()** remains nearly constant time.

Figure 1 shows a bitmap representation of a 512×512 bits random number generated with **bbs.py**. Each black pixel is a 1, and each white pixel is a 0. We use it to visually assess whether the generated random number shows repeated patterns or any discernible bias.

Figures 3 and 4 are a demonstration of why Electronic Code Book mode is semantically insecure for AES encryption. Figure 4 was created by encrypting the body of the original image in PPM format, and then concatenating the unencrypted header back to the encrypted body to create a valid image. Figure 5 was generated using the same protocol but with AES in Cipher Block Chaining mode.




Figure 1: Bitmap representation of a generated pseudo-random number




Figure 2: Modular exponentiation speed comparison

Initiator	Public key equivalence	Shared private key equivalence
A	✓	✓
B	✓	✓

Table 1: Comparing IETF test data to our implementation

Table 2 shows the equivalence between NIST test data⁸ and data generated with our implementation using pycryptodome's AES cipher. The keys, ciphertexts and plaintexts used can be found in the `files/AES_test_data` folder.

Key Length	Plaintext equivalence	Ciphertext equivalence
128 bits	✓	✓
192 bits	✓	✓
256 bits	✓	✓

Table 2: Comparing NIST test data to our implementation

4 Discussion

With a 2048-bit group, the exponent could be as large as $(2^{2048} - 1) \sim 10^{616}$, which would take roughly 10^{170} years to compute with naive exponentiation. The `pow()` method is not only convenient, it is required for the program to complete.

The key generation using Diffie-Hellman with a 2048-bit group is pretty much unbreakable with the current technology as it is estimated that an academic team can break a 768-bit prime and that a nation-state can break a 1024-bit prime⁹. This is arguably the most secure part of our protocol. Thanks to the `pow()` method, this tool can efficiently produce 2048-bit keys, thus generating a $2^{2048} - 1$ key space, which is computationally infeasible to brute-force. Given a generous 1ms for each attempt, it would require $\frac{2^{2048}}{3.154 \times 10^{10}} \sim 10^{626}$ CPU-years to complete an exhaustive search. Nevertheless, having to use 2048-bit keys is a constraint on the efficiency of the algorithm.




Figure 3: Original image




Figure 4: Encrypted image using AES-ECB



Figure 5: Encrypted image using AES-CBC

According to a 2016 NIST report¹⁰, the same security could be achieved using 255-bit keys with Elliptic Curve Diffie-Hellman. Such an implementation would be a bit more complex, but more time and memory efficient.

Concerning BBS, it is hard to conclude anything without running an extensive suite of statistical randomness tests. Visually speaking, the bitmap looks pretty random, but this is only a basic intuition and humans are known to be pretty bad at identifying randomness¹¹. In real world applications, BBS should not be used as a general purpose CSPRNG for multiple reasons. We use it here because we only need to generate a few hundred bits. First of all, BBS is excruciatingly slow, as we can only safely extract a few bits from each iteration. Secondly, the proof of security which comes with this algorithm is only significant for more than 5000-bits moduli^{12 13}. A efficient alternative would be to use AES in Counter mode or in Output Feedback mode, in order to convert a block cipher into a stream cipher generating pseudo-random bits.

Finally, we can see why the Electronic Code Book is an insecure mode of operation for the AES cipher. In Figure 4, the white part of the original image is always encrypted to the same binary data, and is thus recognizable. The rest of the image, because of its irregularities, is pretty much unrecognizable. This is why we assume that for short text messages, ECB is secure enough to encrypt data without an attacker being able to identify repeating patterns. If Alice and Bob wanted to send encrypted images they would have to use another mode, such as Cipher Block Chaining, as demonstrated in Figure 5 where the original image structure cannot be inferred from the encrypted image.

5 Conclusion

We saw how a secure communication protocol could be implemented, so that two actors could communicate confidential data over an insecure channel. Using public key cryptography to exchange keys, key strength reinforcement with pseudo-random number generators and symmetric ciphers, we were able to create an end-to-end encryption scheme. However, this protocol is only a proof of concept, and should be improved to be useful in real-world applications. Transitioning to Elliptic Curve Diffie-Hellman would be a good start for the first primitive. Then, we should use a better CSPRNG than Blum Blum Shub, like AES-CTR or ChaCha20, which are more widely used (TLS, SSL, SSH) and accepted as efficient cryptographically strong pseudo-random generators. Finally, AES should be used in any other mode than ECB. It would then be necessary to communicate IV or nonces as well as data, but this would not be more complicated and it would offer higher security.

References

- [1] Whitfield Diffie and Martin E. Hellman. New directions in cryptography, 1976.
- [2] S. L. Graham, R. L. Rivest, and Ralph C. Merkle. Secure communications over insecure channels, 1978.

- [3] L. Blum, M. Blum, and M. Shub. A simple unpredictable pseudo-random number generator. *SIAM Journal on Computing*, 15(2):364–383, May 1986. doi: 10.1137/0215025. URL <https://doi.org/10.1137/0215025>.
- [4] Joan Daemen and Vincent Rijmen. Aes proposal: Rijndael, 1999.
- [5] Matt Lepinski and Stephen Kent. Additional Diffie-Hellman Groups for Use with IETF Standards. RFC 5114, January 2008.
- [6] Andrew Rukhin, Juan Soto, James Nechvatal, Miles Smid, and Elaine Barker. A statistical test suite for random and pseudorandom number generators for cryptographic applications, 2001.
- [7] Helder Eijs and open source contributors. Pycryptodome’s documentation. <https://www.pycryptodome.org/en/latest/index.html>, 2019.
- [8] Computer Security Division, Information Technology Laboratory, National Institute of Standards, Technology, and Department of Commerce. Example values - cryptographic standards and guidelines. <https://csrc.nist.gov/projects/cryptographic-standards-and-guidelines/example-values>.
- [9] David Adrian, Karthikeyan Bhargavan, Zakir Durumeric, Pierrick Gaudry, Matthew Green, J. Alex Halderman, Nadia Heninger, Drew Springall, Emmanuel Thomé, Luke Valenta, Benjamin VanderSloot, Eric Wustrow, Santiago Zanella-Béguelin, and Paul Zimmermann. Imperfect forward secrecy: How Diffie-Hellman fails in practice. In *22nd ACM Conference on Computer and Communications Security*, October 2015.
- [10] Elaine Barker. Recommendation for key management part 1: General. Technical report, January 2016. URL <https://doi.org/10.6028/nist.sp.800-57pt1r4>.
- [11] Paul A. Warren, Umberto Gostoli, George D. Farmer, Wael El-Deredy, and Ulrike Hahn. A re-examination of “bias” in human randomness perception. *Journal of Experimental Psychology: Human Perception and Performance*, 44(5):663–680, May 2018. doi: 10.1037/xhp0000462. URL <https://doi.org/10.1037/xhp0000462>.
- [12] Sanjit Chatterjee, Alfred Menezes, and Palash Sarkar. Another look at tightness. In Ali Miri and Serge Vaudenay, editors, *Selected Areas in Cryptography*, pages 293–319, Berlin, Heidelberg, 2012. Springer Berlin Heidelberg. ISBN 978-3-642-28496-0.
- [13] Blum blum shub vs. aes-ctr or other csprngs - cryptography stack exchange. <https://crypto.stackexchange.com/a/3456/54184>.